

Bibliography

- [1] S. J. Brams, “Fair Division in Dispute Resolution,” *The Oxford Handbook of Public Choice, Volume 1*. Oxford University Press, 2019. doi: 10.1093/oxfordhb/9780190469733.013.44.
- [2] M. Aleksandrov, H. Aziz, S. Gaspers, and T. Walsh, “Online Fair Division: analysing a Food Bank problem.” [Online]. Available: <https://arxiv.org/abs/1502.07571>
- [3] M. Mertzaniadis, A. Psomas, and P. Verma, “Automating Food Drop: The Power of Two Choices for Dynamic and Fair Food Allocation.” [Online]. Available: <https://arxiv.org/abs/2406.06363>
- [4] I. Caragiannis, D. Kurokawa, H. Moulin, A. D. Procaccia, N. Shah, and J. Wang, “The Unreasonable Fairness of Maximum Nash Welfare,” *ACM Trans. Econ. Comput.*, vol. 7, no. 3, Sep. 2019, doi: 10.1145/3355902.
- [5] H. Steinhaus, “The problem of fair division,” *Econometrica*, vol. 16, pp. 101–104, 1948.
- [6] H. Aziz and S. Mackenzie, “A Discrete and Bounded Envy-Free Cake Cutting Protocol for Any Number of Agents.” [Online]. Available: <https://arxiv.org/abs/1604.03655>
- [7] R. J. Lipton, E. Markakis, E. Mossel, and A. Saberi, “On approximately fair allocations of indivisible goods,” in *Proceedings of the 5th ACM Conference on Electronic Commerce*, 2004, pp. 125–131.
- [8] B. R. Chaudhury, J. Garg, and K. Mehlhorn, “EFX Exists for Three Agents,” *J. ACM*, vol. 71, no. 1, Feb. 2024, doi: 10.1145/3616009.
- [9] H. Akrami, N. Alon, B. R. Chaudhury, J. Garg, K. Mehlhorn, and R. Mehta, “EFX allocations: Simplifications and improvements,” *arXiv preprint arXiv:2205.07638*, 2022.
- [10] B. Plaut and T. Roughgarden, “Almost envy-freeness with general valuations,” *SIAM Journal on Discrete Mathematics*, vol. 34, no. 2, pp. 1039–1068, 2020.
- [11] E. Budish, “The Combinatorial Assignment Problem: Approximate Competitive Equilibrium from Equal Incomes,” *Journal of Political Economy*, vol. 119, no. 6, pp. 1061–1103, 2011, doi: 10.1086/664613.
- [12] H. Chan, J. Chen, B. Li, and X. Wu, “Maximin-Aware Allocations of Indivisible Goods.” [Online]. Available: <https://arxiv.org/abs/1905.09969>
- [13] G. Amanatidis, E. Markakis, and A. Ntotos, “Multiple birds with one stone: Beating $1/2$ for EFX and GMMS via envy cycle elimination,” *Theoretical Computer Science*, vol. 841, pp. 94–109, 2020, doi: <https://doi.org/10.1016/j.tcs.2020.07.006>.
- [14] G. Amanatidis, A. Filos-Ratsikas, and A. Sgouritsa, “Pushing the Frontier on Approximate EFX Allocations.” [Online]. Available: <https://arxiv.org/abs/2406.12413>

- [15] M. Babaioff, T. Ezra, and U. Feige, “Fair and truthful mechanisms for dichotomous valuations,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, 2021, pp. 5119–5126.
- [16] D. Kurokawa, A. Procaccia, and J. Wang, “When Can the Maximin Share Guarantee Be Guaranteed?,” *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 30, no. 1, Feb. 2016, doi: 10.1609/aaai.v30i1.10041.
- [17] G. J. Woeginger, “A polynomial-time approximation scheme for maximizing the minimum machine completion time,” *Operations Research Letters*, vol. 20, no. 4, pp. 149–154, 1997.
- [18] J. Garg, P. McGlaughlin, and S. Taki, “Approximating Maximin Share Allocations,” in *2nd Symposium on Simplicity in Algorithms (SOSA 2019)*, J. T. Fineman and M. Mitzenmacher, Eds., in Open Access Series in Informatics (OASIs), vol. 69. Dagstuhl, Germany: Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 2019, pp. 1–11. doi: 10.4230/OASIs.SOSA.2019.20.
- [19] A. D. Procaccia and J. Wang, “Fair enough: guaranteeing approximate maximin shares,” in *Proceedings of the Fifteenth ACM Conference on Economics and Computation*, in EC '14. Palo Alto, California, USA: Association for Computing Machinery, 2014, pp. 675–692. doi: 10.1145/2600057.2602835.
- [20] J. Garg and S. Taki, “An improved approximation algorithm for maximin shares,” in *Proceedings of the 21st ACM Conference on Economics and Computation*, 2020, pp. 379–380.
- [21] H. Akrami, J. Garg, E. Sharma, and S. Taki, “Simplification and improvement of MMS approximation,” *arXiv preprint arXiv:2303.16788*, 2023.
- [22] H. Akrami and J. Garg, “Breaking the $3/4$ barrier for approximate maximin share,” in *Proceedings of the 2024 Annual ACM-SIAM Symposium on Discrete Algorithms (SODA)*, 2024, pp. 74–91.
- [23] U. Feige, A. Sapir, and L. Tauber, “A Tight Negative Example for MMS Fair Allocations,” in *Web and Internet Economics*, M. Feldman, H. Fu, and I. Talgam-Cohen, Eds., Cham: Springer International Publishing, 2022, pp. 355–372.
- [24] G. Amanatidis, G. Birmpas, A. Filos-Ratsikas, A. Hollender, and A. A. Voudouris, “Maximum Nash welfare and other stories about EFX,” *Theoretical Computer Science*, vol. 863, pp. 69–85, 2021a, doi: <https://doi.org/10.1016/j.tcs.2021.02.020>.
- [25] M. Roos and J. Rothe, “Complexity of social welfare optimization in multiagent resource allocation,” in *Proceedings of the 9th International Conference on Autonomous Agents and Multiagent Systems: volume 1-Volume 1*, 2010, pp. 641–648.
- [26] S. Barman, S. K. Krishnamurthy, and R. Vaish, “Greedy algorithms for maximizing Nash social welfare,” *arXiv preprint arXiv:1801.09046*, 2018a.

- [27] G. Amanatidis, G. Birmpas, A. Filos-Ratsikas, A. Hollender, and A. A. Voudouris, “Maximum Nash welfare and other stories about EFX,” *Theoretical Computer Science*, vol. 863, pp. 69–85, 2021b, doi: <https://doi.org/10.1016/j.tcs.2021.02.020>.
- [28] S. Barman, S. K. Krishnamurthy, and R. Vaish, “Finding fair and efficient allocations,” in *Proceedings of the 2018 ACM Conference on Economics and Computation*, 2018b, pp. 557–574.
- [29] R. Freeman, S. Sikdar, R. Vaish, and L. Xia, “Equitable allocations of indivisible goods,” *arXiv preprint arXiv:1905.10656*, 2019.
- [30] S. Barman, U. Bhaskar, Y. Pandit, and S. Pyne, “Nearly Equitable Allocations beyond Additivity and Monotonicity,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, 2024, pp. 9494–9501.
- [31] A. Farhadi, M. Hajiaghayi, M. Latifian, M. Seddighin, and H. Yami, “Almost envy-freeness, envy-rank, and nash social welfare matchings,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, 2021, pp. 5355–5362.
- [32] C. Basteck and L. Ehlers, “Strategy-proof and envy-free random assignment,” *Journal of Economic Theory*, vol. 209, p. 105618, 2023.
- [33] S. Barman, A. Biswas, S. Krishnamurthy, and Y. Narahari, “Groupwise maximin fair allocation of indivisible goods,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, 2018.
- [34] B. R. Chaudhury, T. Kavitha, K. Mehlhorn, and A. Sgouritsa, “A little charity guarantees almost envy-freeness,” *SIAM Journal on Computing*, vol. 50, no. 4, pp. 1336–1358, 2021.
- [35] D. Kurokawa, “Fair division in game theoretic settings,” 2017.
- [36] H. Aziz, S. Bouveret, I. Caragiannis, I. Giagkousi, and J. Lang, “Knowledge, fairness, and social constraints,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, 2018.
- [37] I. Caragiannis, J. Garg, N. Rathi, E. Sharma, and G. Varricchio, “New Fairness Concepts for Allocating Indivisible Items,” 2023.
- [38] H. Akrami and N. Rathi, “Epistemic EFX Allocations Exist for Monotone Valuations.” [Online]. Available: <https://arxiv.org/abs/2405.14463>
- [39] B. Berger, A. Cohen, M. Feldman, and A. Fiat, “(Almost Full) EFX Exists for Four Agents (and Beyond).” [Online]. Available: <https://arxiv.org/abs/2102.10654>
- [40] H. Akrami, R. Rezvan, and M. Seddighin, “An EF2X allocation protocol for restricted additive valuations,” *arXiv preprint arXiv:2202.13676*, 2022.
- [41] M. Aleksandrov and T. Walsh, “Online fair division: A survey,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, 2020, pp. 13557–13562.

- [42] G. Benade, A. M. Kazachkov, A. D. Procaccia, and C.-A. Psomas, “How to make envy vanish over time,” in *Proceedings of the 2018 ACM Conference on Economics and Computation*, 2018, pp. 593–610.
- [43] M. Aleksandrov and T. Walsh, “Strategy-proofness, envy-freeness and pareto efficiency in online fair division with additive utilities,” in *PRICAI 2019: Trends in Artificial Intelligence: 16th Pacific Rim International Conference on Artificial Intelligence, Cuvu, Yanuca Island, Fiji, August 26–30, 2019, Proceedings, Part I 16*, 2019, pp. 527–541.
- [44] V. Christoforidis and C. Santorinaios, “On The Pursuit of EFX for Chores: Non-Existence and Approximations.” [Online]. Available: <https://arxiv.org/abs/2406.10752>
- [45] S. Barman, V. V. Narayan, and P. Verma, “Fair Chore Division under Binary Supermodular Costs.” [Online]. Available: <https://arxiv.org/abs/2302.11530>
- [46] S. Ebadian, D. Peters, and N. Shah, “How to fairly allocate easy and difficult chores,” *arXiv preprint arXiv:2110.11285*, 2021.
- [47] I. Caragiannis, C. Kaklamanis, P. Kanellopoulos, and M. Kyropoulou, “On low-envy truthful allocations,” in *Algorithmic Decision Theory: First International Conference, ADT 2009, Venice, Italy, October 20-23, 2009. Proceedings 1*, 2009, pp. 111–119.
- [48] G. Amanatidis, G. Birmpas, and E. Markakis, “On truthful mechanisms for maximin share allocations,” *arXiv preprint arXiv:1605.04026*, 2016.
- [49] G. Amanatidis, G. Birmpas, G. Christodoulou, and E. Markakis, “Truthful allocation mechanisms without payments: Characterization and implications on fairness,” in *Proceedings of the 2017 ACM Conference on Economics and Computation*, 2017, pp. 545–562.
- [50] D. Halpern, A. D. Procaccia, A. Psomas, and N. Shah, “Fair Division with Binary Valuations: One Rule to Rule Them All.” [Online]. Available: <https://arxiv.org/abs/2007.06073>
- [51] V. Conitzer, R. Freeman, and N. Shah, “Fair public decision making,” in *Proceedings of the 2017 ACM Conference on Economics and Computation*, 2017, pp. 629–646.
- [52] H. Aziz, H. Moulin, and F. Sandomirskiy, “A polynomial-time algorithm for computing a Pareto optimal and almost proportional allocation,” *Operations Research Letters*, vol. 48, no. 5, pp. 573–578, 2020.
- [53] A. Baklanov, P. Garimidi, V. Gkatzelis, and D. Schoepflin, “PROPM allocations of indivisible goods to multiple agents,” *arXiv preprint arXiv:2105.11348*, 2021.
- [54] S. Boodaghians, B. R. Chaudhury, and R. Mehta, “Polynomial Time Algorithms to Find an Approximate Competitive Equilibrium for Chores,” in *Proceedings of the 2022 Annual ACM-SIAM Symposium on Discrete Algorithms (SODA)*, 2022, pp. 2285–2302.

- [55] B. R. Chaudhury, J. Garg, P. McGlaughlin, and R. Mehta, “On the existence of competitive equilibrium with chores,” in *13th Innovations in Theoretical Computer Science Conference (ITCS 2022)*, 2022.
- [56] A. Mas-Colell, M. Whinston, and J. Green, *Microeconomic Theory*. in Microeconomic Theory. Oxford University Press, 1995. [Online]. Available: <https://books.google.com/books?id=KGtegVXqD8wC>
- [57] H. Hosseini, J. Kavner, S. Sikdar, R. Vaish, and L. Xia, “Hide, Not Seek: Perceived Fairness in Envy-Free Allocations of Indivisible Goods.” [Online]. Available: <https://arxiv.org/abs/2212.04574>
- [58] N. Dupuis-Roy and F. Gosselin, “An empirical evaluation of fair-division algorithms,” in *Proceedings of the annual meeting of the cognitive science society*, 2009.

Introduction

Fair division is a problem with roots in the Bible and in Greek mythology. When Abraham and Lot must divide the land of Canaan fairly among themselves, one divides the land into two parts (which they perceive to be equally valuable), and the other chooses their favorite piece. This problem has proven to be useful in practice. It has been used for international dispute resolution (Brams 2019), distribution of food for food banks, and estate divisions (Aleksandrov et al. 2015; Mertzanidis, Psomas, and Verma 2024; Caragiannis et al. 2019). This was formalized by (Steinhaus 1948) for both divisible continuous and indivisible discrete goods. Since then, several promising results have been published in both settings, with the existence of a bounded-time envy-free cake cutting procedure being resolved by (Aziz and Mackenzie 2017). However, the fair allocation of a finite set of indivisible goods to individuals “fairly” remains a fundamental open question within computer science and economics. Indeed, many interpretations of fairness exist. In this paper, we survey some common notions of fairness, their computability, and open questions within the field of Fair Division. We also present a small result connecting two notions of fairness (MNX and EQX).

Preliminaries

Let \mathcal{A} be a set of agents with \mathcal{A} and \mathcal{G} be a set of goods with \mathcal{G} . Each agent $a \in \mathcal{A}$ possesses a valuation function $v_a: 2^{\mathcal{G}} \rightarrow \mathbb{R}$, representing their value of each bundle of goods within \mathcal{G} . It is assumed that for each agent $a \in \mathcal{A}$ that $v_a(\emptyset) = 0$. For simplicity, we write v_a as v_a for individual items $g \in \mathcal{G}$. The valuation function is usually taken to be additive, meaning that an agent’s value of a bundle is equal to the sum of the agent’s values of the individual goods. Formally, $v_a(S \cup T) = v_a(S) + v_a(T)$ for all $S, T \subseteq \mathcal{G}$. For more general contexts, the valuation function is taken as monotone, i.e., that $v_a(S) \leq v_a(S \cup T)$ for all $S, T \subseteq \mathcal{G}$. In this survey, we focus on the additive context, however several of the results hold for monotone valuations as well.

An allocation α is a partition \mathcal{G} into disjoint sets and assign each partition to an agent $a \in \mathcal{A}$. The subset of \mathcal{G} assigned to agent

$\mathbf{html.frame}$ (called $\mathbf{html.frame}$'s bundle) is denoted $\mathbf{html.frame}$, and each agent $\mathbf{html.frame}$'s valuation of their own bundle is $\mathbf{html.frame}$'s "utility". The goal of fair division is to take a tuple $\mathbf{html.frame}$ and to compute a "fair" allocation efficiently. To this end, there are several notions of fairness. Computing fair allocations efficiently is not always known to be possible, or is hard, so the valuation functions are restricted. We primarily focus on two settings, that of identical valuations where $\mathbf{html.frame}$, and binary valuations where $\mathbf{html.frame}$. In these contexts, several notions become easier to compute.

Envy

Arguably, the gold standard of fair division is envy freeness (EF), where no agent would like another's bundle more than their own. Formally, $\mathbf{html.frame}$. If this is not the case, (i.e. $\mathbf{html.frame}$), $\mathbf{html.frame}$ is said to "envy" "j". It is clear that for discrete goods, this does not always exist, as seen in the following counterexample:

Proof: Take two agents 1 and 2 and one positively valued good. Without loss of generality, allocate the good to agent 1. Agent 2 envies agent 1, as $\mathbf{html.frame}$ and $\mathbf{html.frame}$.

Worse still, computing if an arbitrary instance admits an EF allocation is NP-complete (Lipton et al. 2004).

Thus, when allocating a finite set of goods, EF is relaxed to envy-freeness up to X (EFX) by allowing the removal of an arbitrary good from another agent's bundle such that $\mathbf{html.frame}$ for all $\mathbf{html.frame}$.

For two agents, EFX can be computed using a variation of the "cut-and-choose" procedure. The first agent divides $\mathbf{html.frame}$ into two bundles, $\mathbf{html.frame}$ and $\mathbf{html.frame}$ such that they value each of them as near to $\mathbf{html.frame}$ as possible. Then, the second agent chooses the bundle which gives her higher utility, breaking ties arbitrarily. The second agent clearly does not envy the first, as they received the bundle with the highest valuation in their eyes. If the first agent prefers the bundle they received, we are done. Otherwise, the first agent cannot envy the second after the removal of a good, as otherwise this brings the valuation of the bundles closer to $\mathbf{html.frame}$, which would contradict that the first agent selected as near to $\mathbf{html.frame}$ as they could initially. For three agents, (Chaudhury, Garg, and Mehlhorn 2024) showed that EFX can be computed as well in pseudo-polynomial time. (Akrami et al. 2022) later gave a simpler algorithm for this using casework to iteratively improve the allocation.

Thus, we come to our first open question.

Open Problem 0.1: Do EFX allocations always exist? If so, can they be computed efficiently?

The leximin solution is one which maximizes the minimum utility, then maximizing the second least utility, and so on. The leximin++ is the leximin solution, which, after maximizing the value of the bundle with minimum utility, maximizes size of the bundle. Then, maximizes the size of the second least agent's utility subject to maximizing the size of the second bundle and so on.

For identical additive valuations, an EFX allocation can be computed by taking the leximin++ + maximal allocation. (Plaut and Roughgarden 2020)

Proof: For contradiction, let x be a leximin++ allocation which is not EFX. Then there is a pair of agents i, j with $v_i(x_i) > v_i(x_j)$ and $v_j(x_j) > v_j(x_i)$. This means that the minimum utility of the bundle can be increased by giving that good x_i to j and x_j to i , contradicting that x is a leximin++ allocation. \square

Although computing EFX for general additive valuations remains an open problem, several relaxations of EFX have been proposed which are easier to compute.

One natural easy-to-compute notion relaxation of EFX is EF1. Rather than allowing the removal of an arbitrary good, this is relaxed so that there must exist a good which can each be removed in order to resolve envy. In other words, for a good g , $v_i(x_i - g) \geq v_i(x_j - g)$. This notion was first formally defined by (Budish 2011). Note that this is equivalent to removing the good which i values the most in j 's bundle to enforce envy-freeness up to one good, with i and j . For general valuations, EF1 can be computed efficiently using simple polynomial time algorithms.

One such well-known algorithm is the round robin algorithm, presented below: Informally, the algorithm iterate through agents in a fixed order, giving them their most valued good at each step (breaking ties arbitrarily) until no goods remain.

Round-Robin

Input: $(v_i)_{i \in [n]}$

Output: An EF1 allocation X

```

1  $i \leftarrow 1$ 
2  $P \leftarrow G$ 
3 while  $|P| > 0$ 
4    $g \leftarrow \arg \max_{g \in P} v_i(g)$ 
5    $x_i \leftarrow x_i \cup \{g\}$ 
6    $P \leftarrow P \setminus \{g\}$ 
7  $i \leftarrow i + 1$ 
  
```

(Caragiannis et al. 2019) Round-Robin produces an EF1 allocation

Proof: Take two agents `html.frame` and `html.frame`. After the execution of fixed-order round-robin. First, we note that an agent `html.frame` cannot envy only value another agent `html.frame` with `(html.frame)`. This is trivial for `html.frame`, so let us take `html.frame`. Here, in each iteration of the loop, `html.frame` chooses their favorite good after `html.frame`. Thus, `html.frame` believes they received a weakly more valuable good than `html.frame`. Thus, we get `html.frame`

Note that in each iteration of the loop, for `html.frame`, `html.frame` receives weakly more than what `html.frame` will receive in the following iteration of the loop according to `html.frame`. We then only need to consider the first iteration of the loop, as envy cannot come from later in the execution of the algorithm. Since `html.frame` may value `html.frame` for `html.frame` by the item allocated in the first iteration of the loop, then the removal of the item will remove `html.frame`'s envy of `html.frame`. Thus, the allocation is EF1. `html.frame`

Another common algorithm for computing an EF1 is Envy-Cycle Elimination. Define the envy-graph `html.frame`, with `html.frame` and a directed edge between agents, `html.frame` iff `html.frame`. The algorithm works as follows. Find an arbitrary unenvied agent `html.frame` and give them their most valued good. If no such agent exists, then there must be a cycle `html.frame` where `html.frame` envies `html.frame` (as otherwise there would have to exist a sink, as it would be a DAG). To remove or “eliminate” this cycle (then for agent `html.frame`, give them the bundle of agent `html.frame`). Until an unenvied agent is found, repeatedly remove directed cycles. Then, give the agent their most favorite unallocated good.

Envy-Cycle Elimination

Input: `html.frame`

Output: An EF1 and `html.frame`-EFX allocation `html.frame`

```

1 html.frame
2 for html.frame
3   | html.frame
4 for html.frame
5   | while html.frame unenvied agent # Perform envy-cycle elimination
6     | Find an envy cycle html.frame
7     | html.frame
8   | let html.frame be an unenvied agent # Find an unenvied agent html.frame
9   | let html.frame # Give html.frame their favorite good.
10  | html.frame
11  | html.frame
```

(Lipton et al. 2004) Envy-Cycle Elimination produces an EF1 allocation

Proof: Before the algorithm begins, the allocation is trivially EF (and EFX), as no good has been allocated and thus all bundles are valued at 0 by all agents. In each iteration of the loop, the algorithm allocates goods only to unenvied agents. This implies that an agent may only envy another by a single good. This then implies that the algorithm is EF1 after each iteration of the loop, as the removal of a good from a bundle causing envy for any agent who envies that bundle will resolve the envy. \square

html.frame-EFX

Another relation of EFX is html.frame-EFX, where $u_i(S_i) \geq u_i(S_j)$ for all i, j .

Open Problem 0.2: What is the greatest html.frame for which an html.frame-EFX allocation exists?

It is important to note that if the best possible html.frame, then EFX allocations can always exist, but otherwise, this question becomes meaningful as well.

(Plaut and Roughgarden 2020) first attempt to compute approximate-EFX allocations by slightly modifying the Envy-Cycle elimination algorithm to compute html.frame-EFX. (Chan et al. 2019) show that Envy-Cycle Elimination by itself can compute a html.frame-EFX allocation, if ties are broken between unenvied agents in favor of those with empty bundles.

(Chan et al. 2019) Envy-Cycle Elimination produces an html.frame-EFX allocation

Proof: Take two distinct agents, i and j . Assume that $u_i(S_i) \geq u_i(S_j)$ for all i, j . This statement holds trivially for $i = j$, so assume that $i \neq j$.

Let g be the good last allocated to i . Since ties are broken in favor of empty bundles when picking an agent to allocate to, i , as otherwise j would have been allocated to i rather than j . Because $u_i(S_i) \geq u_i(S_j)$, there must have been a round before g was allocated to i during which i was able to pick an item. Although g might be swapped with other agents' bundles, $u_i(S_i)$ is non-decreasing. Therefore $u_i(S_i) \geq u_i(S_j)$. Additionally, by the design of the algorithm, $u_j(S_j) \geq u_j(S_i)$. Together, these imply that $u_i(S_i) \geq u_i(S_j)$ and $u_j(S_j) \geq u_j(S_i)$.

Thus, the allocation is html.frame-EFX. \square

(Amanatidis, Markakis, and Ntokos 2020) improve this to a html.frame-EFX allocation. Most recently, (Amanatidis, Filos-Ratsikas, and Sgouritsa 2024) compute a html.frame-EFX allocation for general valuations.

Pareto Optimality

Take two allocations x and y . x is said to “Pareto Dominate” y iff $x_i \geq y_i$ for all i , and for at least one agent, $x_i > y_i$. This induces a partial ordering on the set of all feasible allocations. The maximal element(s) are said to be “Pareto Optimal”.

(Plaut and Roughgarden 2020) PO and EFX do not coincide for all fair division instances if zero marginal utility is allowed.

Proof:

	a	b	c
agent 1	x_1	x_2	x_3
agent 2	y_1	y_2	y_3

x and y for a PO allocation, x_1 must be given to agent 1. Similarly, x_2 must be allocated to agent 1 as $x_2 > y_2$ and x_3 .

Assume without loss of generality x_1 is allocated to agent 1. Then x_2 .

Like Pareto dominating allocations, Lorenz dominating allocations induce a partial ordering over the set of allocations. Here, take two vectors of utilities which have been sorted x and y . x is said to “Lorenz dominate” y iff for every k , $\sum_{i=1}^k x_i \geq \sum_{i=1}^k y_i$ (Babaioff, Ezra, and Feige 2021).

MMS

A Maximin-Share Fair allocation is one maximizing the “worst-off” agent according to each x_i . In other words, all agents x_i , x_i .

Unfortunately MMS may not always exist, (Kurokawa, Procaccia, and Wang 2016), although it exists with high probability in general. Additionally, MMS is hard to compute (Woeginger 1997). Thus, several recent papers try to approximate it so that x_i , with x_i . This is known as an x_i -MMS allocation.

One common technique for producing x_i -MMS allocations is known as the bag-filling algorithm. First, the MMS values of all agents are normalized such that x_i for all agents x_i . Maintain a “bag”, and keep adding arbitrary goods until an agent x_i values the bag at least x_i , breaking ties arbitrarily. Once such an agent is found, add all the goods in the bag to x_i and remove them from x_i . When a single agent remains, give her the remaining unallocated goods.

Bag-Filling

Input: x_i

Bag-Filling

Output: A html.frame-MMS allocation X

```
1 html.frame
2 html.frame
3 while html.frame such that html.frame
4   | html.frame # Give html.frame their favorite good.
5   | html.frame
6   | html.frame
7 while html.frame
8   | while html.frame with html.frame
9   |   | html.frame arbitrary html.frame
10  |   | html.frame
11  | Take html.frame where html.frame
12  | html.frame
13  | html.frame
14 Let html.frame be the remaining agent in html.frame
15 html.frame
```

(Garg, McGlaughlin, and Taki 2019) Removing an agent and removing a single good from a problem instance produces results in a weakly greater MMS value. In other words, $u_i(X) \geq u_i(Y)$ for an arbitrary $html.frame$.

Proof: Before the removal of the agent and good, $html.frame$. Suppose w.l.o.g that $html.frame$, and reallocate all other $html.frame$ to the other agents $html.frame$. This is clearly an allocation $html.frame$, denoted by $html.frame$. Since each agent has weakly greater utility in the new allocation, $html.frame$. $html.frame$

(Garg, McGlaughlin, and Taki 2019) The bag-filling algorithm produces a $html.frame$ -MMS allocation.

Proof: In each round of the algorithm, the bag is allocated with an agent with value at least $html.frame$ for the bag. Note that the value of the other agents for the given bag is smaller than 1. Since, before the last good is added to the bag, the value of the bag for all agents is smaller than $html.frame$ and greater than $html.frame$ for at least one agent after. Thus, when an agent takes away a bag in some round, the remaining $html.frame$ agents have total value at least $html.frame$ of the remaining bundle. This can be repeated until 1 agent remains, who will then receive their entire MMS value. $html.frame$

(Procaccia and Wang 2014) first showed the existence of $html.frame$ -MMS approximation algorithms. (Garg, McGlaughlin, and Taki 2019) present a simple algorithm for computing such allocations by modifying the bag-filling algorithm to use low, medium, and highly

valued goods to compute a html.frame-MMS allocation. Several works attempt to (Garg and Taki 2020) compute html.frame-MMS allocations, achieving up to an- html.frame-MMS allocation (Akrami et al. 2023), but with no improvement on the constant approximation factor. Very recently, (Akrami and Garg 2024) improve this to a html.frame-MMS .

Despite the progress made in computing approximations, the approximation of MMS is incomplete. (Feige, Sapir, and Tauber 2022) showed a counterexample, showing that MMS allocations cannot always be approximated higher than html.frame , even for 3 agents.

Open Problem 0.3: What is the highest possible html.frame for which html.frame-MMS allocations exist?

Open Problem 0.4: Does a stronger approximation bound exist than html.frame ?

Maximum Nash Welfare

Maximum Nash Welfare (MNW) is a more “global” notion of welfare. Namely, it considers the utility of each agent rather than agents’ value over others’. This first maximizes the number of agents receiving positive utility, then maximizes the geometric mean of utilities. Let html.frame represent the set of all partitions of a set S into k elements. The MNW allocation is given by html.frame .

In the notion of the geometric mean, under identical additive valuations, this is maximized when all agents with positive utilities have valuations which are as similar as possible (after maximizing the agents with positive utilities). This is quite useful, as MNW provides strong guarantees for other notions of fairness.

(Caragiannis et al. 2019) A MNW allocation is PO and EF1

Proof: MNW implies PO, as a Pareto dominating allocation would have a higher product of utilities. It is left to show that MNW html.frame EF1.

Assume for contradiction that MNW html.frame EF1. This implies that there exists html.frame where html.frame .

Proving EF1 is slightly more involved. Assume for contradiction that there exists two agents html.frame and html.frame such that html.frame for every html.frame . We must show that there is another allocation html.frame with Nash Welfare larger than that of html.frame .

Because html.frame envies html.frame , there exists a html.frame . Let $\text{html.frame} = \text{html.frame}$ for all html.frame . Next, move the html.frame from html.frame ’s bundle to html.frame ’s. In other words, html.frame and html.frame .

We must show that html.frame

To begin, we know that html.frame , which implies that

html.frame , since html.frame envies html.frame even after the removal of an item html.frame in html.frame . Therefore,

html.frame

html.frame

which, from the previous inequality, $\text{html.frame} \text{html.frame} \text{html.frame}$

(Caragiannis et al. 2019) also showed that a MNW allocation is guaranteed to be html.frame -MMS.

These guarantees however do not extend much further. (Amanatidis et al. 2021a) show that MNW does not give an EFX allocation in general, MNW is NP-hard even for identical additive valuations (Roos and Rothe 2010), although under binary and bi-valued (where html.frame) instances it can be computed in polynomial time (Barman, Krishnamurthy, and Vaish 2018a). For 3-value instances, where html.frame , the problem of computing a MNW allocation is NP-complete (Amanatidis et al. 2021b). Thus, due to its hardness to compute, we seek to guarantee these properties without computing them through MNW.

Open Problem 0.5: Can an EF1 and PO allocation be computed in polynomial time?

As computing MNW allocations is APX-hard, thus prior research has sought to approximate MNW allocations for its desiderata. Can an EF1 and PO allocation be computed in polynomial time? For bounded valuations, (Barman, Krishnamurthy, and Vaish 2018b) develop a polynomial time algorithm, and in pseudo-polynomial time for general valuations.

Open Problem 0.6: What is the best possible html.frame for which html.frame -MNW allocations exist?

EQX

Another notion of fairness based on utilities is EQX. An equitable (EQ) allocation is one in which all agent utilities are roughly the same. html.frame , however like EFX, this clearly does not exist. This motivates a similar relaxation, by allowing the removal of a good html.frame . EQX is well studied in the literature for both goods and for chores, and unlike EFX, is fairly easy to compute.

For identical valuations, EFX coincides with EQX, both of which can be computed efficiently. For strictly positive valuations, an EQX + PO allocation always exists, and an EQ1 + PO allocation can be found in pseudo-polynomial time (Freeman et al. 2019).

(Barman et al. 2024) shows that for all chores or all goods, EQX exists even for monotone allocations, and can be computed in polynomial time, although this is not true for arbitrary valuations (unless $P=NP$).

$\mathbf{html.frame}$ under identical additive valuations.

Proof: Assume for contradiction that MNW does not imply EQX under identical valuations. Start with an MNW allocation $\mathbf{html.frame}$. This means that there is an agent from which $\mathbf{html.frame}$. Otherwise, if $\mathbf{html.frame}$ EQX holds trivially.

If $\mathbf{html.frame}$, then it suffices to show that moving a good from $\mathbf{html.frame}$ to $\mathbf{html.frame}$ will result in a larger geometric mean.

$\mathbf{html.frame}$

$\mathbf{html.frame}$

$\mathbf{html.frame}$

By the assumption,

$\mathbf{html.frame}$

and since valuations are additive $\mathbf{html.frame}$, so $\mathbf{html.frame}$. Thus we have derived a contradiction, as $\mathbf{html.frame}$ is increased by moving the good $\mathbf{html.frame.html.frame}$

Between EFX and EF1

EFL (Envy-Freeness up to a Less Preferred Good) is meant to be “in-between” EFX and EF1. For agents $\mathbf{html.frame}$ and $\mathbf{html.frame}$, rather than allowing the removal of an arbitrary good from $\mathbf{html.frame}$, it only allows the removal of goods $\mathbf{html.frame}$ which are less valuable to $\mathbf{html.frame}$ than $\mathbf{html.frame}$. In other words $\mathbf{html.frame}$ for $\mathbf{html.frame}$ and $\mathbf{html.frame}$. This can be computed in polynomial time using envy-cycle elimination.

Randomness

Another line research focuses on randomness in fairness allocations.

EFR

(Farhadi et al. 2021) introduced the notion of envy-freeness up to a random good (EFR), where the expected envy of $\mathbf{html.frame}$ towards an agent $\mathbf{html.frame}$ can be eliminated. (Farhadi et al. 2021) also show that a $\mathbf{html.frame}$ -EFR allocation can be computed in polynomial time. However, such a mechanism cannot be strategyproof and envy free even in expectation (Basteck and Ehlers 2023).

Improving the Maximin Share

GMMS

EFX can be connected to MMS through the stronger notion of (Groupwise MMS) GMMS.

(Barman et al. 2018) define an allocation as GMMS as if for all subsets of $html.frame$ and their bundle, their allocation is also MMS. $html.frame$. It is clear that because MMS allocations do not exist, GMMS allocation allocations do not always exist.

PMMS

PMMS is a slightly weaker form of GMMS, in which $html.frame$. In other words, for all $html.frame$, $html.frame$.

A slightly easier notion is $html.frame$ -GMMS, where $html.frame$. (Barman et al. 2018) show that $html.frame$ -GMMS allocations always exist, and can be computed efficiently. (Amanatidis, Markakis, and Ntokos 2020) show that improved this to a $html.frame$ approximation. The currently highest approximation is to $html.frame$ -(Chaudhury et al. 2021).

A similar line of work has attempted to approximate PMMS, with the highest known approximation as an $html.frame$ 0.781-PMMS allocation (Kurokawa 2017).

Open Problem 0.7: What is the highest $html.frame$ for which $html.frame$ -GMMS allocations exist?

GMMS connects MMS (and its related notions) to EF and its relaxations through the hierarchy:

$html.frame$

(Barman et al. 2018) GMMS strictly generalizes MMS and PMMS.

Proof: It is trivially true that $html.frame$ and that $html.frame$, as MMS must hold for all subsets for $html.frame$ for $html.frame$ respectively by the definition of GMMS. For showing that $html.frame$, we provide the example given by (Barman et al. 2018).

Take an instance with $html.frame$ agents and $html.frame$ goods, where $html.frame$. This means that $html.frame$, since some agents must receive nothing. $html.frame$, and $html.frame$. For $html.frame$ and all other agents $html.frame$.

Take an allocation $html.frame$, with all other agents receiving $html.frame$. This is both MMS and PMMS, but is not GMMS. $html.frame$

Other EFX Notions

Epistemic EFX

(Aziz et al. 2018) proposed the notion of Epistemic Envy-Freeness where for each agent i , the allocation can be made envy-free for i by redistributing i 's bundle among i . (Caragiannis et al. 2023) propose the slightly weaker notion of Epistemic EFX, defined similarly, but only requiring that the allocation can be made envy-free up to i for i . (Akrami and Rathi 2024) recently showed that EEFX allocations always exist and can be computed in polynomial time.

Open Problem 0.8: Do Epistemic Envy-Free allocations exist in general?

Charity

Although the existence of EX allocations is open, it always exists for a subset of goods. For example, trivially, the removal of all goods yields an EFX allocation. The meaningful form of this is the idea of “charity”, where a set of goods remain unallocated (i.e. “donated” to charity), and no agent should envy charity. (Chaudhury et al. 2021) give a polynomial time algorithm which donates at most n goods to charity. (Berger et al. 2021) improve this to donating n goods, and give an algorithm which donates only 1 good in the case of 4 agents.

Open Problem 0.9: Does EFX exist while donating a sublinear number of goods to charity?

Note that if an EFX allocation with 0 goods may be donated to charity exists, this is simply an EFX allocation.

(Akrami, Rezvan, and Seddighin 2022) propose a similar notion called EF k X, in which the removal of k goods from another agent's bundle rather than just one (i.e. EF1X is the same as EFX). They show the existence of EF2X allocations for restricted additive valuations (where $v_i(\{x, y\}) \leq v_i(x) + v_i(y)$). note that an EFX allocation with charity can be turned into an EF2X allocation by giving the n goods to n unique agents.

Different Setting

Online Fair Division

While in the traditional model of fair division, items are assumed to be standard, agents and/or goods may arrive in an online manner. (Aleksandrov and Walsh 2020) survey online fair division in depth. Here, either agents, goods, or both arrive online at times t . One approach to online fairness mechanisms is to asymptotically allocate goods such that the envy-per-good $e_i(t)/g(t)$ tends to 0, where $e_i(t)$ is the maximum envy

one agent has. (Benade et al. 2018) show that randomly allocating goods removes the expected envy. More interestingly, it is impossible to do better than such a mechanism, which is optimal up to a logarithmic factor. Online fair division however is not unilaterally easy to compute. (Aleksandrov and Walsh 2019) show that an ex-ante envy free (i.e. envy-free in expectation) and PO allocation may not always exist.

(Aleksandrov and Walsh 2019) A simultaneously ex-ante envy-free and PO allocation does not always exist.

Proof:

Agent	html.frame	html.frame
1	1	2
2	2	1

By the definition of ex-ante envy-freeness, each agent must have equal probability of receiving each item (i.e. html.frame), however to ensure ex-ante Pareto optimality, each agent must only be given their favorite item. html.frame

Chores and Mixed Manna

When all goods are chores, i.e. html.frame, the problem of fair division becomes much more difficult, and leads to several open problems. EFX for chores does not exist, and indeed is NP-complete to compute (Christoforidis and Santorinaios 2024). Should such an allocation exist, it is mutually exclusive from PO (Barman, Narayan, and Verma 2023). It also cannot have a finite approximation. Still, under the restricted context of bi-valued and binary valuations, EF1 and PO allocations exist (Ebadian, Peters, and Shah 2021; Barman, Narayan, and Verma 2023).

The division of goods html.frame with arbitrary valuations, where html.frame have shown several barriers for computation. For EQX allocations, (Barman et al. 2024) showed that EQX allocations are not guaranteed to exist, even under additive normalized valuations, and that determining if such an allocation exists is strongly NP-hard.

Incentives

Consider a setting with strategic agents. Here, an agent may misreport their valuations of goods in order to receive a better allocation. A *strategyproof* allocation is one in which such misreporting is not beneficial for such an agent.

Strategyproof mechanisms may not always exist while fulfilling the desiderata. (Caragiannis et al. 2009) show that in the context of strategic agents, an EFX allocation may not exist. (Amanatidis, Birmpas, and Markakis 2016) showed that no truthful mechanism could output an html.frame-MMS allocation for html.frame. Indeed, overall truthful mechanisms (Amanatidis et al. 2017) are incompatible allocations for general valuations.

For some restricted contexts, however, this can be done. Namely for binary valuations, computing an MNW and Lorenz-dominating allocation (and thus trivially a PO and EF1)

allocation are strategyproof. (Halpern et al. 2020; Babaioff, Ezra, and Feige 2021) both develop algorithms to compute such allocations. They further show that under binary valuations, and even submodular valuations, a truthful mechanism exists to compute EFX allocations.

Proportionality

One of the earliest notions of fairness is that of proportionality (PROP), where each good values their allocation at least $\frac{1}{n}$ of its value. Clearly, like envy-freeness, it cannot be guaranteed in the case of discrete goods, so it is necessary to weaken it slightly by allowing the addition of goods. In proportionality up to 1 good (PROP1), there exists a good g such that $u_i(g) \geq \frac{1}{n}$ for all i . Similarly, in proportionality up to x (PROPx), all goods g satisfy $u_i(g) \geq \frac{x}{n}$ for all i .

(Conitzer, Freeman, and Shah 2017) showed that a PROP1 and PO allocation always exists and can be computed in polynomial time.

Still proportionality can be disappointing, as (Aziz, Moulin, and Sandomirskiy 2020) showed that even for simple cases, a PROPx allocation may not exist.

A PROPx allocation may not always exist.

Proof: Take 3 agents with identical valuations, and 5 goods. For all agents $u_i(g) = 1$ for $g \in \{1, 2, 3\}$ and $u_i(g) = 0$ for $g \in \{4, 5\}$. No allocation admits a PROPx allocation. \square

An intermediary notion, PROP_m has been shown to always exist and be computable in polynomial time. (Baklanov et al. 2021) defines the maximin good as g_m is the largest among the smallest goods in other agent's bundles. If adding this good's value to $u_i(g_m)$'s utility is greater than or equal to $\frac{1}{n}$, i.e. $u_i(g_m) + v_i(g_m) \geq \frac{1}{n}$, then an allocation is PROP_m.

Fisher Markets

Although Nash Welfare is compelling, for its simultaneous EF1 + PO, it is also NP-hard to compute. One line of work attempting to simultaneously achieve fairness and efficiency tries to use the computation of market equilibria of fisher markets. This has also been used to compute approximate MNW, as mentioned previously. In the fisher market, goods are given prices p_i , agents are given earnings e_i , and monotone valuations, as defined above. These goods may be fractionally allocated with linear valuations, i.e. $u_i(g) = p_i \cdot x_i$ where x_i is the fraction of g given to i .

A fisher market equilibrium is a fractional allocation in which

1. All goods g allocated, i.e. that $\sum_i x_i(g) = 1$
2. For all agents i , $u_i(g) = e_i$
3. All agents only receive goods with maximum bang per buck, i.e. for all g , if $u_i(g) > 0$, then $\frac{u_i(g)}{p_i} \geq \frac{u_i(h)}{p_h}$ for all h with $u_i(h) > 0$

When $e_i = \frac{e}{n}$ for all agents i , this is called a Competitive Equilibrium for Equal Incomes (CEEI).

Open Problem 0.10: Can an exact CEEI be computed in polynomial time?

Although this remains an open question, prior work seeks to compute an approximate html.frame -CEEI allocation. In such an allocation, agents may spend html.frame for some html.frame .

Some work has been dedicated to computing approximations for CEEI. (Boodaghians, Chaudhury, and Mehta 2022) computes a html.frame -CEEI allocation with running time proportional to html.frame . (Chaudhury et al. 2022) gives an algorithm to compute a html.frame -CEEI in polynomial time for html.frame .

A fractionally Pareto optimal allocation (fPO) is a fractional allocation which is not Pareto-dominated by another fractional allocation. It is clear that an fPO allocation is PO, but not necessarily vice versa.

CEEI is compelling in computing fPO. (Mas-Colell, Whinston, and Green 1995) show that a CEEI is also fPO, and thus PO.

Empirical Results

We can see the practical usefulness of these results from prior experimental results. For discrete fair division, (Hosseini et al. 2023) collected a set of items, and computed an EF1 allocation among the agents. They then allowed agents to swap, and found that agents preferred to swap their bundles with other agents in html.frame of allocations.

Another common approach to allowing measurement of goods is to give each agent some number of “points”, and allow them to distribute the points among the goods based on their preference, which can still maintain the fairness properties of MNW and EF1. In Spliddit, for 1000 points and 10 agents, an MNW (EF1 and PO) allocation is computed in less than 3 seconds (Caragiannis et al. 2019; Dupuis-Roy and Gosselin 2009).

Discussion

In this survey, we discussed some of the most common notions of fairness. Although there are some further related notions of fairness, these are the most well-studied among the literature. We also showed that, under identical additive valuations, html.frame .

Related to fair division, within economics and computation (EC), recent research has also studied computing market equilibria, beyond the Fisher Market. Other works in EC focus more on strategyproofness, studying voting and consensus procedures.